Healthcare AI/ML Engineer Roadmap

Action Kit

Build predictive models, clinical NLP systems, and medical imaging AI that directly impact patient care and clinical workflows.

Difficulty	Very High
Timeline	12 to 18 months
Last Updated	2026-03-24

This action kit contains your 90-day execution plan, portfolio project templates, resume bullet formulas, interview preparation questions, and resource checklist.

Built by Ehoneah Obed | thetransmutation.substack.com | quiz.ehoneahobed.com

Section 1: 90-Day Execution Plan

This plan maps directly to your learning path phases. Each month has specific deliverables and checkpoints to keep you on track.

Month 1: Foundation

Hours: 200 to 250

- Month 1 to 2: Python Fundamentals (syntax, data types, functions, OOP)
- pandas DataFrames (the tool you will use 90% of the time)
- Basic SQL for querying relational databases
- Month 2 to 3: Statistics for Data Science (probability, hypothesis testing, significance)
- Correlation vs. causation (critical in healthcare)
- Bias and variance, overfitting, cross-validation
- Month 3 to 4: Machine Learning Fundamentals
- Supervised vs. unsupervised learning
- Model evaluation metrics (accuracy, precision, recall, AUC-ROC)
- Common algorithms: logistic regression, decision trees, random forests

Checkpoint: Write a Python script that reads a CSV, calculates summary statistics, and filters data using a public healthcare dataset. Build a complete ML pipeline: load, split, train multiple models, evaluate, explain.

Month 2: Technical Depth

Hours: 300 to 350

- Month 5 to 6: Deep Learning and Neural Networks
- Neural network architecture, CNNs for image data, RNNs/LSTMs for sequential data
- Transformer architectures (important for medical imaging and clinical NLP)
- Month 7 to 8: Healthcare Data Formats and Clinical NLP
- FHIR and HL7 standards
- EHR data structure and how it differs from clean datasets
- Clinical NLP: extracting entities from clinical notes
- Month 8 to 9: ML Engineering and Production Systems
- Version control (Git), experiment tracking, model serving
- Data pipelines, ETL, monitoring and retraining

Checkpoint: Train a CNN on a medical imaging dataset (CheXpert). Parse FHIR data and extract patient information, OR build an NLP model that extracts medications from clinical notes. Add version control, parameter tracking, evaluation metrics to your deep learning model using MLflow.

Month 3: Specialization

Hours: 200 to 300

- Track A: Medical Imaging AI (CNNs, object detection, segmentation; datasets: CheXpert, LUNA16, BraTS)
- Track B: Clinical NLP and EHR Data (NER, clinical coding, cohort identification; datasets: MIMIC-III/IV, I2B2)
- Track C: Predictive Analytics and Clinical Risk Models (time-series, survival analysis, causal inference; datasets: MIMIC, Framingham)
- Track D: Drug Discovery and Pharmaceutical AI (molecular prediction, pharmacogenomics; tools: DeepChem, RDKit; datasets: PubChem, ChEMBL)

Checkpoint: Polished end-to-end project in your chosen specialization, documented on GitHub.

Section 2: Portfolio Project Templates

Each project below includes a dataset, tools, timeline, and your clinical advantage. Complete at least 2 to 3 of these for a competitive portfolio.

Project 1: Hospital Readmission Prediction

Build complete ML pipeline predicting 30-day readmission from patient data, including feature engineering, model selection, evaluation, and clinical validation.

Dataset	UCI Diabetes Readmission
URL	https://archive.ics.uci.edu/dataset/468/diabetes+130+us+hospitals+for+years+1999+2008
Tools	Python, pandas, scikit-learn
Time Estimate	20 to 30 hours

Clinical Advantage: You recognize that missing data patterns have clinical meaning and validate predictions against clinical intuition

Project 2: Medical Imaging Classification

CNN classifying chest X-rays to detect pneumonia with documented model architecture, training process, evaluation metrics, and clinical interpretation.

Dataset	CheXpert
URL	https://stanfordmlgroup.github.io/competitions/chexpert/
Tools	Python, TensorFlow, PyTorch
Time Estimate	30 to 40 hours

Clinical Advantage: You understand why model interpretability matters for radiologists and limitations of trained models

Project 3: Clinical NLP: Medication Extraction

NLP model extracting medication names and dosages from clinical notes with entity recognition, evaluated on real clinical text.

Dataset	MIMIC-III
URL	https://mimic.physionet.org/
Tools	spaCy, scispaCy, Hugging Face transformers
Time Estimate	35 to 45 hours

Clinical Advantage: You recognize abbreviations and understand why context matters in clinical note analysis

Project 4: Time-Series Patient Deterioration Prediction

LSTM or transformer model predicting sepsis or acute decompensation from continuous vital signs with temporal pattern analysis.

Dataset	MIMIC time-series data or PhysioNet SEPSIS dataset
URL	https://mimic.physionet.org/
Tools	PyTorch, TensorFlow, LSTM or attention models
Time Estimate	40 to 50 hours

Clinical Advantage: You understand clinical deterioration patterns and the value of early warning systems

Project 5: End-to-End Specialization Project

Polished project in your chosen track (imaging, NLP, predictive analytics, or drug discovery) with fairness analysis, model explainability, and clinical deployment considerations.

Tools	Full ML engineering stack appropriate to specialization
Time Estimate	50 to 70 hours

Clinical Advantage: Deep domain expertise guides every design decision from problem framing to model validation

Section 3: Resume Bullet Formulas

Use these formulas to translate your clinical experience into language that resonates with hiring managers in health data roles. Replace bracketed text with your specifics.

Nursing Background

• Leveraged [clinical skill: Pattern recognition and clinical reasoning] to deliver [tech outcome: Model validation and error analysis: You integrate vital signs, patient appearance, context into assessment. In ML, you instinctively ask whether the model's error profile differs across patient subgroups.], resulting in [quantified impact].

• Leveraged [clinical skill: HIPAA and privacy compliance] to deliver [tech outcome: Data governance and de-identification workflows: You already navigate patient privacy daily and understand why removing a name is not de-identification.], resulting in [quantified impact].

• Leveraged [clinical skill: Documentation under pressure] to deliver [tech outcome: Model documentation and explainability: Writing clear clinical notes translates to writing model cards and explaining outputs to non-technical clinicians.], resulting in [quantified impact].

• Leveraged [clinical skill: Workflow optimization in resource-constrained settings] to deliver [tech outcome: Pragmatic engineering and MVP thinking: Healthcare AI is full of 95% accurate models that never get used because they require data that does not exist in that hospital's workflow.], resulting in [quantified impact].

Pharmacy Background

• Leveraged [clinical skill: Medication data analysis and safety] to deliver [tech outcome: Drug-drug interaction prediction and adverse event detection: Your mental checklist of contraindications and interactions becomes the foundation for building models that catch medication errors.], resulting in [quantified impact].

• Leveraged [clinical skill: Pharmacokinetics and pharmacodynamics] to deliver [tech outcome: Personalized medicine and precision dosing models: Understanding how drug concentrations change based on individual physiology translates to precision dosing models in oncology and critical care.], resulting in [quantified impact].

• Leveraged [clinical skill: Regulatory knowledge and FDA compliance] to deliver [tech outcome: Regulated ML and FDA Software as a Medical Device (SaMD): FDA's 2021 guidance on AI/ML in software as medical device is where your regulatory literacy becomes a direct career advantage.], resulting in [quantified impact].

• Leveraged [clinical skill: Clinical literature evaluation] to deliver [tech outcome: Domain-specific AI research translation: You already synthesize evidence and evaluate study design limitations.], resulting in [quantified impact].

Other Clinical Background

• Leveraged [clinical skill: Complex diagnostic reasoning under uncertainty] to deliver [tech outcome: Medical imaging AI and diagnostic decision support: Understanding what 'normal variation' looks like and why a model trained on high-quality radiology scans fails in rural clinics with older equipment.], resulting in [quantified impact].

• Leveraged [clinical skill: Patient risk stratification] to deliver [tech outcome: Predictive analytics and clinical risk models: You already think about risk factors, protective factors, and how they interact.], resulting in [quantified impact].

• Leveraged [clinical skill: Longitudinal patient care] to deliver [tech outcome: Time-series ML and temporal pattern recognition: Your intuition for how conditions evolve over time is a superpower for sequential patient data models.], resulting in [quantified impact].

• Leveraged [clinical skill: Interprofessional collaboration] to deliver [tech outcome: Healthcare AI implementation and change management: The graveyard of healthcare AI is full of technically perfect models that never got adopted.], resulting in [quantified impact].

Pro Tip: Always quantify. Instead of 'analyzed data,' write 'analyzed 50,000+ claims records to identify $2.3M in recoverable overpayments.'

Section 4: Interview Preparation Questions

Practice these role-specific questions. For each, prepare a 2-minute answer that highlights both your technical skills and clinical background.

Technical Questions

1. How would you approach building a clinical NLP model to extract diagnoses from unstructured clinical notes?
2. Describe the steps you would take to validate an ML model for clinical use, including bias detection.
3. How do you handle class imbalance in a dataset of rare adverse events?
4. Walk me through your approach to deploying an ML model into a production clinical workflow.
5. How would you ensure a predictive model complies with FDA SaMD regulations?

Behavioral Questions

1. Tell me about a time your clinical experience helped you identify a bias or error in a model that others missed.
2. How do you explain model limitations and uncertainty to clinical stakeholders?
3. Describe how you balance model accuracy with interpretability for clinical end-users.
4. How do you handle pushback from clinicians who do not trust AI recommendations?

Scenario Questions

1. A deployed sepsis prediction model starts generating false positives at a higher rate than in testing. What do you do?
2. A hospital wants to use your model on a patient population very different from your training data. How do you advise them?
3. You discover that your training data contains a systematic bias against a specific demographic group. What is your response?

Section 5: Resource Checklist

Use this checklist to track your progress through all recommended resources.

Certifications

Certification	Signal	Cost	Timeline
AWS Certified Machine Learning, Specialty	High	$300 exam	2 to 3 months study, 35 to 50% of job postings
Google Cloud Professional ML Engineer	High	$200 exam	2 to 3 months study, common at startups and GCP-forward companies
TensorFlow Developer Certificate	Helpful	$100 exam	1 to 2 months study, 15 to 30% of postings
Fast.ai Practical Deep Learning Certificate	Helpful	Free course plus $50 optional cert	N/A
Andrew Ng ML Specialization	Helpful	~$200	N/A

Learning Resources by Phase

Foundation

- [Free] Python for Everybody: https://www.py4e.com/
- [Free] Kaggle Python course: https://www.kaggle.com/learn/python
- [Free] SQLZoo: https://sqlzoo.net/
- [Free] StatQuest with Josh Starmer: https://www.youtube.com/c/joshstarmer
- [Free] Seeing Theory: http://students.brown.edu/seeing-theory/
- [Free] Google ML Crash Course: https://developers.google.com/machine-learning/crash-course
- [Free] Fast.ai Practical Deep Learning: https://course.fast.ai/
- [Free] scikit-learn tutorials: https://scikit-learn.org/stable/getting_started.html
- [Paid] Andrew Ng ML Specialization on Coursera: https://www.coursera.org/specializations/machine-learning-introduction

Technical Depth

- [Free] 3Blue1Brown Neural Networks series: https://www.youtube.com/channel/UCYO_jab_esuFRV4b17AJtAw
- [Free] Fast.ai Part 2: https://course.fast.ai/
- [Free] TensorFlow tutorials: https://www.tensorflow.org/tutorials
- [Paid] DeepLearning.AI Specialization on Coursera: https://www.deeplearning.ai/
- [Free] spaCy NLP tutorials: https://spacy.io/usage
- [Free] scispaCy for medical NLP: https://allenai.github.io/scispacy/

- [Free] ClinicalBERT: https://github.com/kexinhuang12345/ClinicalBERT
- [Free] Made With ML: https://madewithml.com/
- [Free] MLflow documentation: https://mlflow.org/docs/latest/index.html

Specialization

- [Free] Medical Imaging datasets and resources: https://stanfordmlgroup.github.io/competitions/chexpert/
- [Free] MIMIC-III and MIMIC-IV datasets: https://mimic.physionet.org/
- [Free] DeepChem: https://deepchem.io/
- [Free] RDKit: https://www.rdkit.org/

Your First Three Moves

Map your clinical superpower (2 to 3 hours)

- Write a 1 to 2 page document identifying patient safety issues you have noticed that could be caught by algorithms
- Note what data clinicians lack access to
- Describe which clinical decisions take hours but could be automated
- This becomes your personal positioning statement for interviews and project selection

Complete your first Python mini-project (10 to 15 hours)

- Pick one public healthcare dataset (UCI Diabetes Readmission or Breast Cancer Wisconsin)
- Load data, calculate statistics, identify factors associated with outcomes
- Train a basic model
- Push the code to GitHub

Commit to one foundational course (5 to 7 hours per week)

- Google ML Crash Course (free, https://developers.google.com/machine-learning/crash-course) is most recommended
- Alternative: Fast.ai Practical Deep Learning (free, https://course.fast.ai/) if you prefer top-down learning
- Spend 30 minutes to enroll and commit to the schedule

Ready to go deeper? Take the free career quiz at quiz.ehoneahobed.com or subscribe to The Transmutation newsletter at thetransmutation.substack.com